

IE 604: Network Flow Optimization¹

Course Notes, Part II: Shortest Paths

Dr. Jim Ostrowski

University of Tennessee - Knoxville

¹Last updated: February 9, 2022

Section 1


Introduction

Unit 1

The Shortest Path Problem

Notation

- Directed network,
 $G = (N, A)$
- Arcs, $(i, j) \in A$
 - length (cost): c_{ij}
 - max length:
 $C = \max_{ij} \{c_{ij}\}$
- Adjacency list of node i ,
 A_i
- Source, s



Problem Definition

- Length perspective
 - for every nonsource node $i \in N$, find a shortest length directed path from node s to node i
 - the length of a path is the sum of the lengths of the arcs on the path
- Minimum cost flow perspective
 - find a way to send 1 unit of flow as cheaply as possible from node s to every other node $i \in N \setminus \{s\}$ in an uncapacitated network

Linear Programming Formulation

$$\begin{array}{ll}\text{Minimize} & \sum_{(i,j) \in A} c_{ij} x_{ij} \\ \text{subject to} & \sum_{(i,j) \in A} x_{ij} - \sum_{(j,i) \in A} x_{ji} = \begin{cases} n-1 & \text{for } i = s \\ -1 & \forall i \in N \setminus \{s\} \end{cases} \\ & x_{ij} \geq 0.\end{array}$$

Linear Programming Example

How do we obtain the shortest paths from
the optimal solution?


LP: single destination

Linear Programming Dual: single destination

Linear Programming Dual: single destination (cont.)


Assumptions

- 1 All arc lengths are integers
- 2 The network contains a directed path from node s to every other node in the network
- 3 The network does not contain a negative cycle
- 4 Network is directed



Main Types of Shortest Path Problems

- 1 Single-source acyclic graph
- 2 Single-source with non-negative arc lengths on general graph
- 3 Single-source with arbitrary arc lengths on acyclic graph
- 4 Single-source with arbitrary arc lengths on general graph
- 5 All-pairs shortest path problem



Other Problems Related to the Shortest Path Problem


- ① Maximum capacity path problem
- ② Maximum reliability path problem
- ③ Shortest path problem with turn penalties
- ④ Shortest path with an additional constraint
- ⑤ Resource-constrained shortest path problem

Unit 2

Applications

Modeling the Behavior of a Wildfire

- A landscape can be modeled as a **grid graph**
 - nodes represent center points of cells
 - the weights on arcs represent the travel time between two cell midpoints (computed using ellipse model)
 - $R = \frac{b^2 - c^2}{b - c \times \cos(\theta)}$,
($0 \leq \theta < \pi/2$)
 - $R = \frac{b^2 - c^2}{b - c \times \cos(\pi - \theta)}$,
($\pi/2 \leq \theta \leq \pi$)
- The time needed for a fire to spread from cell s to cell t is the **shortest path**




- > Fire spread direction between cells
- > The major fire spread direction in a cell

Modeling the Behavior of a Wildfire (2)

Modeling the Behavior of a Computer Hacker

- Computer security researchers have developed a tool called an **attack graph** to visualize the vulnerability of a computer network
 - nodes represent the **privilege / identity** that the attacker has obtained
 - arcs represent exploits that can be used to escalate the privilege / identity
- We could assign weights to the edges in the attack graph to represent the time needed to
- The minimum time/cost for the hacker to gain root access is the **shortest path** from a leaf node to the root node



Section 2


Introduction to Shortest Path Algorithms

Unit 1

Types of Shortest Path Algorithms

Label Algorithms

- Assign tentative distance labels to nodes at each step
 - labels represent upper bounds on the shortest path to each node




Label-Setting Algorithms

- Procedure
 - designate one label as permanent in each iteration
- Advantages
 - more efficient
- Disadvantages
 - only works for graphs without a negative cycle

Label-Setting Algorithms

Label-Correcting Algorithms

- Procedure
 - consider all labels as temporary until the last step in which they all become permanent
- Advantages
 - work for all classes of problems
 - offer more algorithmic flexibility
- Disadvantages
 - less efficient





Unit 2

How to Store the Shortest Paths

Shortest Path Tree

- Directed-out tree
 - rooted at node s
- Unique path from s to a node i in the tree is a shortest path from s to i
- Why is this true?



Property 4.1

- If the path $s = i_1 - i_2 - \cdots - i_h = k$ is a shortest path from node s to node k , then for every $q = 2, 3, \dots, h - 1$, the subpath $s = i_1 - i_2 - \cdots - i_q$ is a shortest path from the source node to node i_q

Proof.

(By contradiction)

Property 4.2

Let $d(i)$ be the shortest path from s to i

- A directed path P from s to k is a shortest path if and only if $d(j) = d(i) + c_{ij}$ for all $(i, j) \in P$

Proof.

Property 4.2 (Cont.)

Why Does the Tree Contain the Shortest Paths?

- 1 The network contains a finite number of paths; thus, there is a shortest path to every node
- 2 Then, by Property 4.2, we can always find a shortest path from s to an node such that

$$d(j) = d(i) + c_{ij} \quad \forall (i, j) \in P \quad (1)$$


- 3 Perform breadth-first search on the network using the arcs that satisfy (1), creating a breadth-first tree

Unit 3


Acyclic Networks

Shortest Path on Acyclic Networks


- A simple “sweeping” algorithm is optimal
 - number nodes according to their topological ordering
 - set distance label $d(s) = 0$ and the others to a very large number
- for i in *TopologicallyOrderedSet*
 - for $(j, i) \in A_i^-$
 - if $d(i) > d(j) + c_{ji}$
then
 $d(i) \leftarrow d(j) + c_{ji}$




Example



Practice



Solution



Why Does It Work?

Proof.

(By induction.)

Why Does It Work? (2)

Worst Case Complexity

- Each arc is “examined” exactly once
 - addition
 - comparison
 - possible assignment
- Thus, the algorithm solves the problem in $O(m)$ time

Section 3

Dijkstra's Algorithm

Unit 1

Introduction

About Dijkstra's Algorithm

- Edsger Dijkstra (1930-2002)
 - conceived his algorithm in 1956
 - published it in 1959
- Label-setting algorithm
- Single-source shortest path problem
- Differences with algorithm for acyclic networks
 - can be used on any graph without negative cycles
 - may examine an arc more than once




Unit 2

The Algorithm

Terminology

- Nodes divided into two groups
 - permanently labeled, S
 - temporarily labeled, \bar{S}
- Label values, $d(i)$
 - distance to a permanently labeled node is the shortest distance to that node
 - distance to temporarily labeled node is an upper bound on the shortest distance to that node
 - all labels are the distances along paths whose internal nodes are permanently labeled

Terminology (2)



Algorithm Summary


- Basic idea
 - start at node s
 - while there remain temporary nodes
 - **node selection:** choose neighbor with smallest temporary label; make this node a permanent node
 - update the temporary labels of all the neighbors of the new permanent node

$$d(i) = \min_{(j,i) \in A} \{d(j) + c_{ji}\}$$


Unit 3

Examples


Example



Practice



Solution



Unit 4

Technical Details

Pseudocode

- Initialize
 - $S := \emptyset; \bar{S} = N$
 - $d(s) = 0$ and $pred(s) = 0$
 - $d(i) := \infty$ for all $i \in N \setminus \{s\}$
- **while** $|S| < n$
 - let $i \in \bar{S}$ be a node for which $d(i) = \min\{d(j) : j \in \bar{S}\}$
 - $S := S \cup \{i\}; \bar{S} := \bar{S} \setminus \{i\}$
 - **for** each $(i, j) \in A_i^+$
 - **if** $d(j) > d(i) + c_{ij}$ **then** $d(j) := d(i) + c_{ij}$ and $pred(j) := i$

Correctness

Summary of Proof

- Proof by mathematical induction
- Induction hypotheses
 - ① the distance label of each node in S is optimal
 - ② the distance label of each node in S is the shortest path length from the source provided that each internal node in the path lies in S

Correctness

Hypothesis #1

Correctness

Hypothesis #1 (2)

Correctness

Hypothesis #2

Worst-Case Complexity

- **Node selection**
 - performed n times
 - each time requires scanning all temporary nodes (up to n nodes)
 - $n + (n - 1) + (n - 2) + \cdots + 1 = O(n^2)$
- **Distance updates**
 - performed $|A_i|$ times for each node
 - Overall, $\sum_{i \in N} |A_i| = m$ times
 - each time requires $O(1)$ time
 - $O(m)$ total time
- **Overall**
 - $O(n^2 + m) = O(n^2)$
 - cannot get better for completely dense graphs

Section 4

Introduction to Label-Correcting Algorithms

Unit 1

Introduction

Motivation

- Label setting algorithms only work for acyclic networks or networks with non-negative costs (i.e., networks without negative cycles)
- In general, we cannot easily solve networks with negative cycles
- We would be willing to accept an algorithm that can do the following:
 - identifies a negative cycle when one exists
 - solves the problem if no negative cycle exists

How They Work

- Generic algorithm
 - reduces the distance label of one node at each iteration by considering only local information
 - all labels are made permanent when the algorithm terminates
- More advanced algorithms use optimality conditions

Optimality Conditions

Necessary conditions (\Leftarrow)

- If the distance labels $d(j)$ are shortest path distances, **then** they must satisfy the following necessary optimality conditions

$$d(j) \leq d(i) + c_{ij} \quad \forall (i, j) \in A \quad (2)$$

- Proof
 - if this inequality is violated for any arc, then we can reduce the shortest path (contradicting the optimality of the distance labels) to node j by setting it to

$$d(j) = d(i) + c_{ij}$$

Optimality Conditions

Sufficient conditions (\implies)

- If each $d(j)$ represents the length of some directed path from s to node j and this solution satisfies the conditions (2), **then** it must be optimal (i.e., the shortest path)
- Proof

Optimality Conditions: Proof (Cont.)

Optimality Conditions

Theorem: Shortest Path Optimality Conditions

Theorem

For every node $j \in N$, let $d(j)$ denote the length of some directed path from the source node to node j . Then the numbers $d(j)$ represent shortest path distances if and only if they satisfy the following shortest path optimality conditions:

$$d(j) \leq d(i) + c_{ij} \quad \forall (i, j) \in A$$

Properties About Reduced Arc Lengths

- Define the **reduced arc length of an arc (i, j)** with respect to the distance labels $d(\cdot)$, as
 - $c_{ij}^d = c_{ij} + d(i) - d(j)$
- Properties
 - For any directed cycle W
 - $\sum_{(i,j) \in W} c_{ij}^d = \sum_{(i,j) \in W} c_{ij}$
 - For any directed path P from node k to node ℓ
 - $\sum_{(i,j) \in P} c_{ij}^d = \sum_{(i,j) \in P} c_{ij} + d(k) - d(\ell)$
 - If $d(\cdot)$ represent shortest path distances, $c_{ij}^d \geq 0$ for all $(i, j) \in A$


Unit 2

The Generic Label-Correcting Algorithm

The Generic Label-Correcting Algorithm

- Assumptions
 - Assumes a graph with no negative cycles
- Procedure
 - Find some arc (i, j) that violate optimality conditions
 - $d(j) \leq d(i) + c_{ij}$
 - Update the distance label for the arc head
 - $d(j) \leftarrow d(i) + c_{ij}$
 - Terminate when all arcs satisfy the optimality conditions


The Generic Label-Correcting Algorithm




Predecessor Graph

- Directed out-tree T rooted at s
 - collection of arcs $(pred(j), j)$
- When a distance update is made, one arc is removed and another added


Example



Practice



Solution



Pseudocode

- Initialize
 - $d(s) := 0$ and $pred(s) := 0$
 - $d(j) := \infty$ for all $j \in N \setminus \{s\}$
- **while** some arc (i, j) satisfies $d(j) > d(i) + c_{ij}$
 - $d(j) := d(i) + c_{ij}$
 - $pred(j) := i$

Worst-Case Complexity

- Assume data are integral
- Let C be the maximum absolute value of arc costs
- For each finite $d(j)$, $-nC \leq d(j) \leq nC$ because each path contains at most $n - 1$ arcs
- Algorithm updates any label $d(j)$ at most $2nC$ times
- Algorithm performs $O(n^2C)$ iterations

Unit 3

Modified Label-Correcting Algorithm

Introduction

- How do we find an arc that violates the optimality condition?
 - scan the arc list
 - how much time does this require?
 - can we do better than a scan?

List of Potentially Violating Arcs

- Maintain a list, *LIST*, of all arcs that *might* violate their optimality conditions
- Steps
 - ① Select some arc $(i, j) \in LIST$ that violates optimality conditions
 - ① $d(j) > d(i) + c_{ij}$
 - ② Remove (i, j) from *LIST*
 - ③ Update the distance label $d(j) = d(i) + c_{ij}$
 - ① this may cause some arcs in A_j to violate their optimality condition
 - ② decreasing $d(j)$ maintains the optimality condition of for all incoming arcs at node j
 - ④ Add arcs in A_j to *LIST*

Modified Label-Correcting Algorithm

How it Works


- Instead of holding arcs in *LIST*, hold nodes with this property:
 - if an arc (i, j) violates the optimality condition, *LIST* must contain node i

Modified Label-Correcting Algorithm


Pseudocode

- Maintain a list, $LIST$, of all **nodes** that *might* violate their optimality conditions
- **while** $LIST \neq \emptyset$
 - ① remove some node $i \in LIST$
 - ② **for** each $(i, j) \in A_i$
 - if $d(j) > d(i) + c_{ij}$
 - $d(j) = d(i) + c_{ij}$
 - $pred(j) = i$
 - if** $j \notin LIST$ **then** add j to $LIST$

Practice



Solution



Worst-Case Complexity

- Each update of $d(j)$
 - scan each arc in A_i^+ : $O(|A_i^+|)$
- Algorithm can update a label at most $O(2nC)$ times
- Overall bound
 - $\sum_{i \in N} (2nC)|A_i^+| = O(nmC)$

Section 5

Special Implementations of the Modified Label-Correcting Algorithm

Unit 1

Introduction

Why Modify It?

- The generic algorithm does not run in polynomial time
 - $O(n^2C)$

Unit 2

$O(nm)$ Implementation

How It Works

- For $n - 1$ iterations:
 - for all $(i, j) \in A$:
 - if $d(j) > d(i) + c_{ij}$, then set $d(j) := d(i) + c_{ij}$

Correctness and Complexity

Theorem

The label-correcting algorithm requires $O(nm)$ time as long as we examine all arcs at every pass

Proof.

(By induction)



Correctness and Complexity: Proof (Cont.)

FIFO Version

- We don't need to consider every $(i, j) \in A$ during every pass
 - if a node's distance label doesn't change during a pass, we can ignore it during the next pass
- Implementation
 - if a node's distance label changes during a pass, add it to a FIFO queue

Unit 3

Detection of Negative Cycles

Modification to Label-Correcting Algorithm

- Recall that:
 - $-nC \leq d(j) \leq nC$
- Thus, if a distance label falls below $-nC$, then there exists a negative cycle

Modification #2 ($\mathcal{O}(n)$)

DetectCycleInPredecessorGraph

Label source node.

some node is unlabeled

choose unlabeled node k

label k

trace predecessor indices starting at node k , assigning label k to all nodes encountered until first unlabeled node ℓ is reached

nodes k and ℓ have the same labels
return certificate that the PG contains a negative cycle

return certificate that PG has no negative cycle

- if PG contains a cycle, then G contains a negative cycle
- apply check after every αn distance label updates (does not add to worst-case complexity)

Section 6

All-Pairs Shortest Paths

Unit 1

Introduction

All Pairs Shortest Path Problem

- Problem statement
 - “Find the shortest paths between all pairs of nodes in a network”
- Assumptions
 - network is strongly connected
 - no negative cycles
- Algorithms
 - ① Repeated shortest path algorithm
 - ② All-pairs label-correcting algorithm
 - ① generic version
 - ② Floyd-Warshall implementation

Unit 2

Repeated Shortest Path Algorithm

How it Works

- If the network contains negative arcs
 - use the FIFO label-correcting algorithm to transform network
- Solve single-source shortest path algorithm n times on the transformed network (each time using a different node as a source)

Transforming Network With Negative Arcs

FIFO Label-Correcting Algorithm ($\mathcal{O}(nm)$)

- Procedure
 - **If** the network contains a negative cycle, the algorithm will detect it
 - **Else**
 - compute shortest distances $d(j)$ from a source s to all other nodes ($\mathcal{O}(nm)$)
 - set the cost of each arc as the reduced cost:
 $c_{ij}^d := c_{ij} + d(i) - d(j)$
- Computing the shortest path distances in the original network
 - let $\sum_{(i,j) \in P_{\ell k}} c_{ij}$ be the shortest path length from k to ℓ on the transformed network
 - the shortest path length on the original network is
 $\sum_{(i,j) \in P_{\ell k}} c_{ij} + d(\ell) - d(k)$ [see property 5.2(c)]

Complexity

- Let $S(n, m, C)$ be the time needed to solve a shortest path problem with non-negative arc lengths
- The FIFO label-correcting algorithm takes $\mathcal{O}(nm)$ time
- Overall complexity
 - $\mathcal{O}(nm + nS(n, m, C)) = \mathcal{O}(nS(n, m, C))$

Unit 3

All Pairs Optimality Conditions

Notation

- Distance label, $d[i, j]$
 - represents the length of some directed walk from node i to node j
 - is an upper bound on the shortest path from i to j
 - if no walk exists, is infinite

Theorem

Theorem

(All-Pairs Shortest Path Optimality Conditions). For every pair of nodes $[i, j] \in N \times N$, let $d[i, j]$ represent the length of some directed path from node i to node j satisfying $d[i, i] = 0$ for all $i \in N$ and $d[i, j] \leq c_{ij}$ for all $(i, j) \in A$. These distances represent all-pairs shortest path distances if and only if they satisfy the following all-pairs shortest path optimality conditions:

$$d[i, j] \leq d[i, k] + d[k, j] \text{ for all nodes } i, j, \text{ and } k$$

Unit 4

All-Pairs Generic Label-Correcting Algorithm

Summary

- Start with some distance labels $d[i, j]$
- Successively update these until they satisfy the optimality conditions

Pseudocode

AllPairsGenericLabelCorrecting

$d[i, j] := \infty$ for all $[i, j] \in N \times N$

$d[i, i] := 0$ for all $i \in N$

$d[i, j] := c_{ij}$ for all $(i, j) \in A$

$d[i, j] > d[i, k] + d[k, j]$ for some nodes i, j, k

$d[i, j] > d[i, k] + d[k, j] \implies d[i, j] := d[i, k] + d[k, j]$

Complexity

- Assume data are integral
- Let C be the maximum absolute value of arc costs
- For each finite $d[i, j]$, $-nC \leq d[i, j] \leq nC$ because each path contains at most $n - 1$ arcs
- Algorithm updates any label $d[i, j]$ at most $2nC$ times
- The graph has n^2 pairs of nodes
- Algorithm performs $O(n^3 C)$ iterations

Unit 5

Floyd-Warshall Algorithm

Key Property

- Let $d^k[i, j]$ represent the length of the shortest path from i to j subject to the condition that the path only uses nodes $1, 2, \dots, k - 1$ as internal nodes
- Property

$$d^{k+1}[i, j] = \min \left\{ d^k[i, j], d^k[i, k] + d^k[k, j] \right\}$$

How it Works

- 1 Compute $d^1[i, j]$ for all $i, j \in N \times N$
- 2 Use key property to compute $d^2[i, j]$ given $d^1[i, j]$ for all $i, j \in N \times N$

$$d^2[i, j] = \min \{ d^1[i, j], d^1[i, k] + d^1[k, j] \}$$

- 3 Continue until $d^{n+1}[i, j]$ is computed for all $i, j \in N \times N$

Algorithm maintains predecessor indices, $pred[i, j]$, which denote the last node prior to j in the shortest path from i to j

Pseudocode

FloydWarshall

$d[i, j] := \infty$ and $pred[i, j] := 0$ for all $[i, j] \in N \times N$

$d[i, i] := 0$ for all $i \in N$

$d[i, j] := c_{ij}$ and $pred[i, j] := i$ for all $(i, j) \in A$

1 to n


$[i, j] \in N \times N$ if $d[i, j] > d[i, k] + d[k, j]$

$d[i, j] := d[i, k] + d[k, j]$
 $pred[i, j] := pred[k, j]$


Complexity

- Major iterations: n (one for each k)
- Computations per iteration
 - $\mathcal{O}(1)$ for each node pair
- Overall complexity
 - $\mathcal{O}(n \times n^2) = \mathcal{O}(n^3)$

Example



Example: Initialization




Distances

	1	2	3	4
1				
2				
3				
4				

Predecessors

	1	2	3	4
1				
2				
3				
4				

Example: 1




Distances

	1	2	3	4
1				
2				
3				
4				

Predecessors

	1	2	3	4
1				
2				
3				
4				

Example: 2




Distances

	1	2	3	4
1				
2				
3				
4				

Predecessors

	1	2	3	4
1				
2				
3				
4				

Example: 3




Distances

	1	2	3	4
1				
2				
3				
4				

Predecessors

	1	2	3	4
1				
2				
3				
4				

Example: 4



Distances

	1	2	3	4
1				
2				
3				
4				

Predecessors

	1	2	3	4
1				
2				
3				
4				

Unit 6

Detection of Negative Cycles

Generic All-Pairs Label-Correcting Algorithm

- 1 If $i = j$, check whether $d[i, i] < 0$
- 2 If $i \neq j$, check whether $d[i, j] < -nC$

Floyd-Warshall Algorithm

- 1 Check if $d[i, i] < 0$ whenever updating $d[i, i]$ for some node i
- 2 Check if predecessor graph contains a negative cycle